

ZZedc System Administrator Guide

Contents

ZZedc System Administrator Guide	2
For Biostatistics Support Teams and IT Administrators	2
Document Information	2
Table of Contents	2
1. Server Maintenance Procedures	2
1.1 Backup Procedures	2
1.2 Log Rotation	4
1.3 System Updates and Patches	4
2. Database Administration	6
2.1 SQLite Maintenance	6
2.2 Performance Tuning	8
2.3 Connection Pool Management	9
3. Troubleshooting Guide	9
3.1 Common Issues and Solutions	9
3.2 Diagnostic Commands Reference	11
3.3 Recovery Procedures	12
4. Monitoring and Alerting	13
4.1 Health Check Endpoints	13
4.2 Metrics Collection	14
4.3 Alert Configuration	15
4.4 Log Monitoring	17
5. Multi-Site Coordination	17
5.1 Multi-Site Architecture	17
5.2 Site Configuration	18
5.3 Site Access Control	19
5.4 Cross-Site Data Queries	19
5.5 Site Synchronization Monitoring	20
6. Security Operations	20
6.1 Security Monitoring	20
6.2 Session Management	21
6.3 Audit Log Protection	22
7. Upgrade and Patch Management	23
7.1 Version Upgrade Procedure	23
7.2 Rollback Procedure	24
7.3 Security Patch Policy	25
Appendix A: Command Reference	26
System Management	26
R Console Commands	26
Appendix B: Monitoring Dashboards	26
Recommended Metrics for Grafana/Prometheus	26
Key Performance Indicators	26
Document History	27

ZZedc System Administrator Guide

For Biostatistics Support Teams and IT Administrators

Document Information

- **Version:** 1.0.0
 - **Date:** December 2025
 - **Audience:** Biostatisticians, System Administrators, DevOps Engineers
 - **Prerequisites:** Linux/Unix administration, R environment management, database concepts
-

Table of Contents

1. Server Maintenance Procedures
 2. Database Administration
 3. Troubleshooting Guide
 4. Monitoring and Alerting
 5. Multi-Site Coordination
 6. Security Operations
 7. Upgrade and Patch Management
-

1. Server Maintenance Procedures

1.1 Backup Procedures

Automated Daily Backups Configure automated backups using the built-in backup system:

```
#!/bin/bash
# /opt/zzedc/scripts/daily_backup.sh

BACKUP_DIR="/var/backups/zzedc"
DATE=$(date +%Y%m%d_%H%M%S)
DB_PATH="/var/data/zzedc/study.db"
RETENTION_DAYS=30

# Create backup directory
mkdir -p "${BACKUP_DIR}/daily"

# Backup database with integrity check
Rscript -e "
library(zzedc)
backup_result <- perform_database_backup(
  db_path = '${DB_PATH}',
  backup_dir = '${BACKUP_DIR}/daily',
  backup_name = 'zzedc_backup_${DATE}',
  verify = TRUE
)
if (!backup_result\$success) stop('Backup failed: ', backup_result\$error)
"

# Backup configuration files
tar -czf "${BACKUP_DIR}/daily/config_${DATE}.tar.gz" \
  /opt/zzedc/config.yml \\\
```

```

/opt/zzedc/.Renvironment

# Remove backups older than retention period
find "${BACKUP_DIR}/daily" -type f -mtime +${RETENTION_DAYS} -delete

# Log completion
echo "$(date): Daily backup completed" >> /var/log/zzedc/backup.log

```

Cron Schedule Configuration

```

# /etc/cron.d/zzedc-backups

# Daily backup at 2:00 AM
0 2 * * * zzedc /opt/zzedc/scripts/daily_backup.sh

# Weekly full backup on Sunday at 3:00 AM
0 3 * * 0 zzedc /opt/zzedc/scripts/weekly_backup.sh

# Monthly archive on 1st at 4:00 AM
0 4 1 * * zzedc /opt/zzedc/scripts/monthly_archive.sh

```

Backup Verification Procedure Run weekly verification of backup integrity:

```

# Verify backup can be restored
verify_backup_integrity <- function(backup_path) {
  temp_db <- tempfile(fileext = ".db")
  on.exit(unlink(temp_db))

  # Attempt restoration to temporary location
  restore_result <- restore_database_backup(
    backup_path = backup_path,
    restore_path = temp_db,
    verify_only = TRUE
  )

  if (!restore_result$success) {
    warning("Backup verification failed: ", restore_result$error)
    return(FALSE)
  }

  # Verify table counts
  con <- DBI::dbConnect(RSQLite::SQLite(), temp_db)
  on.exit(DBI::dbDisconnect(con), add = TRUE)

  tables <- DBI::dbListTables(con)
  expected_tables <- c("subjects", "visits", "form_data", "audit_trail",
    "users", "data_dictionary")

  missing <- setdiff(expected_tables, tables)
  if (length(missing) > 0) {
    warning("Missing tables in backup: ", paste(missing, collapse = ", "))
    return(FALSE)
  }

  message("Backup verification successful")
}

```

```
    TRUE  
}
```

1.2 Log Rotation

Application Log Configuration Create `/etc/logrotate.d/zzedc`:

```
/var/log/zzedc/*.log {  
    daily  
    rotate 14  
    compress  
    delaycompress  
    missingok  
    notifempty  
    create 0640 zzedc zzedc  
    sharedscripts  
    postrotate  
        systemctl reload zzedc 2>/dev/null || true  
    endscript  
}  
  
/var/log/zzedc/audit/*.log {  
    weekly  
    rotate 52  
    compress  
    delaycompress  
    missingok  
    notifempty  
    create 0640 zzedc zzedc  
    # Audit logs require longer retention for compliance  
}
```

Log Categories and Locations

Log Type	Location	Retention	Purpose
Application	/var/log/zzedc/app.log	14 days	General application events
Error	/var/log/zzedc/error.log	30 days	Error tracking
Audit	/var/log/zzedc/audit/audit7ogears	70 years	Regulatory compliance
Access	/var/log/zzedc/access.log	90 days	User access patterns
Security	/var/log/zzedc/security.log	1 year	Security events

1.3 System Updates and Patches

R Package Updates

```
# Monthly package update procedure  
update_zzedc_packages <- function(dry_run = TRUE) {  
    # Check for available updates  
    old_packages <- old.packages()  
  
    if (is.null(old_packages)) {  
        message("All packages are up to date")
```

```

    return(invisible(NULL))
}

# Filter for ZZedc-related packages
zzedc_deps <- c("shiny", "bslib", "RSQLite", "pool", "DT",
               "ggplot2", "dplyr", "digest", "openxlsx")

to_update <- old_packages[rownames(old_packages) %in% zzedc_deps, ]

if (nrow(to_update) == 0) {
  message("No ZZedc dependency updates available")
  return(invisible(NULL))
}

message("Updates available:")
print(to_update[, c("Package", "Installed", "ReposVer")])

if (!dry_run) {
  # Create restore point
  message("Creating package restore point...")
  renv::snapshot()

  # Apply updates
  update.packages(oldPkgs = to_update, ask = FALSE)
  message("Updates applied. Restart ZZedc service to activate.")
}
}

```

System Update Procedure

```

#!/bin/bash
# /opt/zzedc/scripts/system_update.sh

# 1. Create backup before update
/opt/zzedc/scripts/daily_backup.sh

# 2. Stop ZZedc service
sudo systemctl stop zzedc

# 3. Update system packages
sudo apt update && sudo apt upgrade -y

# 4. Update R packages (in maintenance mode)
Rscript -e "source('/opt/zzedc/scripts/update_packages.R')"

# 5. Run database migrations if needed
Rscript -e "zzedc::run_migrations()"

# 6. Verify application integrity
Rscript -e "zzedc::verify_installation()"

# 7. Restart ZZedc service
sudo systemctl start zzedc

```

```

# 8. Verify service health
sleep 10
curl -s http://localhost:3838/health | grep -q "ok" || {
    echo "Health check failed, rolling back..."
    sudo systemctl stop zzedc
    /opt/zzedc/scripts/restore_backup.sh
    sudo systemctl start zzedc
}

echo "Update completed successfully"

```

2. Database Administration

2.1 SQLite Maintenance

Database Vacuum and Optimization Schedule weekly optimization:

```

# Database optimization procedure
optimize_database <- function(db_path, verbose = TRUE) {
    con <- DBI::dbConnect(RSQLite::SQLite(), db_path)
    on.exit(DBI::dbDisconnect(con))

    # Get initial size
    initial_size <- file.size(db_path)

    if (verbose) message("Running integrity check...")
    integrity <- DBI::dbGetQuery(con, "PRAGMA integrity_check")
    if (integrity$integrity_check != "ok") {
        stop("Database integrity check failed: ", integrity$integrity_check)
    }

    if (verbose) message("Analyzing tables...")
    DBI::dbExecute(con, "ANALYZE")

    if (verbose) message("Running VACUUM...")
    DBI::dbExecute(con, "VACUUM")

    if (verbose) message("Optimizing indexes...")
    DBI::dbExecute(con, "REINDEX")

    final_size <- file.size(db_path)
    savings <- initial_size - final_size

    if (verbose) {
        message(sprintf("Optimization complete. Space reclaimed: %.2f MB",
                       savings / 1024^2))
    }

    list(
        initial_size = initial_size,
        final_size = final_size,
        space_reclaimed = savings,
        integrity = "ok"
    )
}

```

```
)  
}
```

Integrity Check Procedures

```
# Comprehensive database health check  
check_database_health <- function(db_path) {  
  con <- DBI::dbConnect(RSQLite::SQLite(), db_path)  
  on.exit(DBI::dbDisconnect(con))  
  
  results <- list()  
  
  # 1. Integrity check  
  results$integrity <- DBI::dbGetQuery(con, "PRAGMA integrity_check")$integrity_check  
  
  # 2. Foreign key check  
  results$foreign_keys <- DBI::dbGetQuery(con, "PRAGMA foreign_key_check")  
  results$fkViolations <- nrow(results$foreign_keys)  
  
  # 3. Table statistics  
  tables <- DBI::dbListTables(con)  
  results$table_stats <- lapply(tables, function(tbl) {  
    count <- DBI::dbGetQuery(con, sprintf("SELECT COUNT(*) as n FROM %s", tbl))$n  
    list(table = tbl, row_count = count)  
  })  
  
  # 4. Index health  
  results$indexes <- DBI::dbGetQuery(con,  
    "SELECT name, tbl_name FROM sqlite_master WHERE type='index'")  
  
  # 5. Database size  
  results$size_bytes <- file.size(db_path)  
  results$size_mb <- results$size_bytes / 1024^2  
  
  # 6. Page statistics  
  results$page_count <- DBI::dbGetQuery(con, "PRAGMA page_count")$page_count  
  results$page_size <- DBI::dbGetQuery(con, "PRAGMA page_size")$page_size  
  results$freelist_count <- DBI::dbGetQuery(con, "PRAGMA freelist_count")$freelist_count  
  
  # Calculate fragmentation  
  used_pages <- results$page_count - results$freelist_count  
  results$fragmentation_pct <- (results$freelist_count / results$page_count) * 100  
  
  # Generate health score  
  results$health_score <- calculate_health_score(results)  
  
  class(results) <- c("zzedc_db_health", "list")  
  results  
}  
  
calculate_health_score <- function(results) {  
  score <- 100
```

```

if (results$integrity != "ok") score <- score - 50
if (results$fk_violations > 0) score <- score - 20
if (results$fragmentation_pct > 20) score <- score - 10
if (results$fragmentation_pct > 40) score <- score - 10

max(0, score)
}

```

2.2 Performance Tuning

SQLite Pragma Configuration

```

# Apply performance pragmas to database connection
configure_db_performance <- function(con) {
  # Write-ahead logging for better concurrency
  DBI::dbExecute(con, "PRAGMA journal_mode = WAL")

  # Increase cache size (negative = KB, positive = pages)
  DBI::dbExecute(con, "PRAGMA cache_size = -64000") # 64 MB cache

  # Synchronous mode: NORMAL balances safety and speed
  DBI::dbExecute(con, "PRAGMA synchronous = NORMAL")

  # Memory-mapped I/O (256 MB)
  DBI::dbExecute(con, "PRAGMA mmap_size = 268435456")

  # Temporary storage in memory
  DBI::dbExecute(con, "PRAGMA temp_store = MEMORY")

  # Enable foreign keys
  DBI::dbExecute(con, "PRAGMA foreign_keys = ON")

  invisible(con)
}

```

Query Performance Analysis

```

# Analyze slow queries
analyze_query_performance <- function(con, query, explain = TRUE) {
  if (explain) {
    # Get query plan
    plan <- DBI::dbGetQuery(con, paste("EXPLAIN QUERY PLAN", query))
    message("Query Plan:")
    print(plan)
  }

  # Time execution
  start <- Sys.time()
  result <- DBI::dbGetQuery(con, query)
  elapsed <- Sys.time() - start

  list(
    rows_returned = nrow(result),
    execution_time = elapsed,

```

```

        query_plan = if (explain) plan else NULL
    )
}

2.3 Connection Pool Management

# Connection pool configuration
create_db_pool <- function(db_path, min_size = 1, max_size = 5) {
  pool::dbPool(
    drv = RSQLite::SQLite(),
    dbname = db_path,
    minSize = min_size,
    maxSize = max_size,
    idleTimeout = 60000, # 60 seconds
    validationInterval = 10000, # 10 seconds
    onCreate = function(con) {
      configure_db_performance(con)
    }
  )
}

# Monitor pool health
monitor_pool_health <- function(pool) {
  list(
    size = pool::poolSize(pool),
    idle = pool::poolIdle(pool),
    active = pool::poolSize(pool) - pool::poolIdle(pool),
    valid = pool::dbIsValid(pool)
  )
}

```

3. Troubleshooting Guide

3.1 Common Issues and Solutions

Issue: Application Fails to Start **Symptoms:** Service fails to start, port already in use, package loading errors.

Diagnostic Steps:

```

# Check service status
sudo systemctl status zzedc

# View recent logs
sudo journalctl -u zzedc -n 50 --no-pager

# Check port availability
sudo lsof -i :3838

# Verify R environment
Rscript -e "sessionInfo()"

```

Solutions:

```
# Kill orphaned processes
```

```

sudo fuser -k 3838/tcp

# Reset application state
sudo systemctl stop zzedc
rm -f /var/run/zzedc/*.pid
sudo systemctl start zzedc

# Reinstall packages if corrupted
Rscript -e "remove.packages('shiny'); install.packages('shiny')"

```

Issue: Database Connection Errors **Symptoms:** “Database is locked”, connection timeouts, read-only errors.

Diagnostic Steps:

```

# Check database file permissions
file.info("/var/data/zzedc/study.db")

# Check for lock files
list.files("/var/data/zzedc", pattern = "\\\\.db-", full.names = TRUE)

# Test basic connectivity
con <- DBI::dbConnect(RSQLite::SQLite(), "/var/data/zzedc/study.db")
DBI::dbGetQuery(con, "SELECT 1")
DBI::dbDisconnect(con)

```

Solutions:

```

# Fix permissions
sudo chown zzedc:zzedc /var/data/zzedc/study.db
chmod 660 /var/data/zzedc/study.db

# Remove stale locks (CAUTION: ensure no active connections)
sudo systemctl stop zzedc
rm -f /var/data/zzedc/study.db-shm /var/data/zzedc/study.db-wal
sudo systemctl start zzedc

# Recover from WAL
Rscript -e "
con <- DBI::dbConnect(RSQLite::SQLite(), '/var/data/zzedc/study.db')
DBI::dbExecute(con, 'PRAGMA wal_checkpoint(TRUNCATE)')
DBI::dbDisconnect(con)
"

```

Issue: Memory Exhaustion **Symptoms:** Application crashes, R session aborts, out of memory errors.

Diagnostic Steps:

```

# Check system memory
free -h

# Check R memory usage
Rscript -e "gc(); print(gc())"

# Monitor real-time memory
watch -n 1 'ps aux | grep -E "R|zzedc" | grep -v grep'

```

Solutions:

```
# Configure R memory limits in .Renviron
# R_MAX_VSIZE=8Gb
# R_MAX_MEM_SIZE=8Gb

# Implement data pagination for large queries
fetch_paginated <- function(con, table, page_size = 1000) {
  total <- DBI::dbGetQuery(con, sprintf("SELECT COUNT(*) FROM %s", table))[[1]]
  pages <- ceiling(total / page_size)

  lapply(seq_len(pages), function(page) {
    offset <- (page - 1) * page_size
    DBI::dbGetQuery(con, sprintf(
      "SELECT * FROM %s LIMIT %d OFFSET %d",
      table, page_size, offset
    ))
  })
}
```

Issue: Slow Performance **Symptoms:** Page load times > 5 seconds, report generation timeouts.

Diagnostic Steps:

```
# Profile application
profvis::profvis({
  # Run problematic operation
})

# Check database query times
system.time({
  DBI::dbGetQuery(con, "SELECT * FROM form_data WHERE visit_id = 1")
})

# Monitor connection pool
monitor_pool_health(db_pool)
```

Solutions:

```
# Add missing indexes
DBI::dbExecute(con, "CREATE INDEX IF NOT EXISTS idx_form_data_visit
                      ON form_data(visit_id)")
DBI::dbExecute(con, "CREATE INDEX IF NOT EXISTS idx_form_data_subject
                      ON form_data(subject_id)")
DBI::dbExecute(con, "CREATE INDEX IF NOT EXISTS idx_audit_timestamp
                      ON audit_trail(timestamp)")

# Enable query caching
options(zzedc.query_cache_size = 100)
options(zzedc.query_cache_ttl = 300) # 5 minutes
```

3.2 Diagnostic Commands Reference

```
# System diagnostics
/opt/zzedc/scripts/diagnostics.sh

# Database diagnostics
```

```
Rscript -e "zzedc::diagnose_database('/var/data/zzedc/study.db')"

# Connection diagnostics
Rscript -e "zzedc::test_connections()

# Full health check
Rscript -e "zzedc::full_health_check()
```

3.3 Recovery Procedures

Database Recovery from Corruption

```
# Attempt database recovery
recover_database <- function(db_path) {
  backup_path <- paste0(db_path, ".corrupt.", format(Sys.time(), "%Y%m%d%H%M%S"))

  # Preserve corrupted database
  file.copy(db_path, backup_path)
  message("Corrupted database preserved at: ", backup_path)

  # Attempt recovery using .dump
  recovery_path <- paste0(db_path, ".recovered")

  system2("sqlite3", args = c(
    db_path,
    sprintf(".output %s.sql", recovery_path),
    ".dump"
  ))

  # Recreate from dump
  if (file.exists(paste0(recovery_path, ".sql"))) {
    unlink(recovery_path)
    system2("sqlite3", args = c(
      recovery_path,
      sprintf(".read %s.sql", recovery_path)
    ))
  }

  # Verify recovery
  con <- DBI::dbConnect(RSQLite::SQLite(), recovery_path)
  integrity <- DBI::dbGetQuery(con, "PRAGMA integrity_check")
  DBI::dbDisconnect(con)

  if (integrity$integrity_check == "ok") {
    file.rename(recovery_path, db_path)
    message("Database recovered successfully")
    return(TRUE)
  }
}

warning("Automatic recovery failed. Manual intervention required.")
FALSE
}
```

Service Recovery Procedure

```

#!/bin/bash
# /opt/zzedc/scripts/emergency_recovery.sh

echo "ZZedc Emergency Recovery Procedure"
echo "====="

# Stop all services
sudo systemctl stop zzedc

# Clear temporary files
rm -rf /tmp/zzedc_*
rm -rf /var/run/zzedc/*

# Reset log files if corrupted
if ! file /var/log/zzedc/app.log | grep -q "text"; then
    mv /var/log/zzedc/app.log /var/log/zzedc/app.log.corrupt
    touch /var/log/zzedc/app.log
    chown zzedc:zzedc /var/log/zzedc/app.log
fi

# Verify database integrity
Rscript -e "
result <- zzedc::check_database_health('/var/data/zzedc/study.db')
if (result\$health_score < 50) {
    stop('Database health critical. Restore from backup required.')
}
"

# Restart services
sudo systemctl start zzedc

# Verify recovery
sleep 5
if curl -s http://localhost:3838/health | grep -q "ok"; then
    echo "Recovery successful"
else
    echo "Recovery failed. Restore from backup."
    /opt/zzedc/scripts/restore_latest_backup.sh
fi

```

4. Monitoring and Alerting

4.1 Health Check Endpoints

Implementing Health Check API The application exposes health check endpoints for monitoring:

```

# Health check endpoint handler (in server.R)
observeEvent(input$health_check, {
  health <- list(
    status = "ok",
    timestamp = Sys.time(),
    version = packageVersion("zzedc"),
    database = check_db_connection(),
    memory = check_memory_usage(),

```

```

        disk = check_disk_space(),
        sessions = get_active_session_count()
    )

    # Return appropriate status code
    if (any(sapply(health[-c(1,2,3)], function(x) x$status == "critical"))){
        health$status <- "critical"
    } else if (any(sapply(health[-c(1,2,3)], function(x) x$status == "warning"))){
        health$status <- "warning"
    }

    health
)

```

Health Check Script for External Monitoring

```

#!/bin/bash
# /opt/zzedc/scripts/health_check.sh

HEALTH_URL="http://localhost:3838/health"
ALERT_EMAIL="admin@example.org"

response=$(curl -s -w "\n%{http_code}" "$HEALTH_URL" 2>/dev/null)
http_code=$(echo "$response" | tail -n 1)
body=$(echo "$response" | head -n -1)

if [ "$http_code" != "200" ]; then
    echo "CRITICAL: ZZedc health check failed (HTTP $http_code)" |
        mail -s "ZZedc Alert: Health Check Failed" "$ALERT_EMAIL"
    exit 2
fi

status=$(echo "$body" | jq -r '.status')

case "$status" in
    "ok")
        echo "OK: ZZedc is healthy"
        exit 0
        ;;
    "warning")
        echo "WARNING: ZZedc has warnings"
        echo "$body" | mail -s "ZZedc Warning" "$ALERT_EMAIL"
        exit 1
        ;;
    "critical")
        echo "CRITICAL: ZZedc is in critical state"
        echo "$body" | mail -s "ZZedc Critical Alert" "$ALERT_EMAIL"
        exit 2
        ;;
esac

```

4.2 Metrics Collection

System Metrics Dashboard

```

# Collect system metrics for monitoring
collect_system_metrics <- function() {
  metrics <- list(

    timestamp = Sys.time(),

    # Application metrics
    active_users = get_active_session_count(),
    requests_per_minute = get_request_rate(),
    avg_response_time_ms = get_avg_response_time(),

    # Database metrics
    db_connections_active = pool::poolSize(db_pool) - pool::poolIdle(db_pool),
    db_connections_idle = pool::poolIdle(db_pool),
    db_size_mb = file.size(db_path) / 1024^2,

    # System metrics
    memory_used_pct = get_memory_usage_pct(),
    disk_used_pct = get_disk_usage_pct(),
    cpu_load_avg = get_cpu_load(),

    # Data metrics
    total_subjects = get_subject_count(),
    total_records = get_record_count(),
    pending_queries = get_pending_query_count()
  )

  # Log to time series
  append_metrics_log(metrics)

  metrics
}

```

4.3 Alert Configuration

Alert Thresholds

```
# /opt/zzedc/config/alerts.yml
```

```

alerts:
  memory:
    warning: 75
    critical: 90

  disk:
    warning: 80
    critical: 95

  response_time_ms:
    warning: 2000
    critical: 5000

  database_connections:
    warning: 8 # out of 10
    critical: 10

```

```

failed_logins:
  window_minutes: 15
  warning: 5
  critical: 10

error_rate:
  window_minutes: 5
  warning: 0.05 # 5%
  critical: 0.10 # 10%

notifications:
  email:
    enabled: true
    recipients:
      - admin@example.org
      - biostat-support@example.org

slack:
  enabled: false
  webhook_url: ""

```

Alert Handler

```

# Process alerts based on thresholds
process_alerts <- function(metrics, config) {
  alerts <- list()

  # Memory alert
  if (metrics$memory_used_pct >= config$alerts$memory$critical) {
    alerts$memory <- list(
      level = "critical",
      message = sprintf("Memory usage critical: %.1f%%", metrics$memory_used_pct)
    )
  } else if (metrics$memory_used_pct >= config$alerts$memory$warning) {
    alerts$memory <- list(
      level = "warning",
      message = sprintf("Memory usage elevated: %.1f%%", metrics$memory_used_pct)
    )
  }

  # Disk alert
  if (metrics$disk_used_pct >= config$alerts$disk$critical) {
    alerts$disk <- list(
      level = "critical",
      message = sprintf("Disk space critical: %.1f%%", metrics$disk_used_pct)
    )
  }

  # Response time alert
  if (metrics$avg_response_time_ms >= config$alerts$response_time_ms$critical) {
    alerts$performance <- list(
      level = "critical",
      message = sprintf("Response time critical: %d ms", metrics$avg_response_time_ms)
    )
  }
}

```

```

}

# Send notifications
if (length(alerts) > 0) {
  send_alert_notifications(alerts, config$notifications)
}

alerts
}

```

4.4 Log Monitoring

Centralized Log Aggregation

```

# /etc/rsyslog.d/zzedc.conf

# Forward ZZedc logs to central syslog server
if $programname == 'zzedc' then @@logserver.example.org:514

# Local logging with structured format
template(name="zzedc-json" type="list") {
  constant(value="{")
  constant(value="\\"timestamp\\":\\\"") property(name="timereported" dateFormat="rfc3339")
  constant(value=",\\\"severity\\\":\\\"") property(name="syslogseverity-text")
  constant(value=",\\\"message\\\":\\\"") property(name="msg" format="json")
  constant(value="\\\"}\\n")
}

if $programname == 'zzedc' then /var/log/zzedc/structured.log;zzedc-json

```

Log Analysis Queries

```

# Find errors in last hour
grep -E "ERROR|CRITICAL" /var/log/zzedc/app.log |
  awk -v cutoff="$(date -d '1 hour ago' '+%Y-%m-%d %H:%M')" '$1" "$2 >= cutoff'

# Count errors by type
grep "ERROR" /var/log/zzedc/app.log |
  sed 's/.*ERROR: //' |
  sort | uniq -c | sort -rn

# Find slow queries (>1 second)
grep "Query time:" /var/log/zzedc/app.log |
  awk -F': ' '$2 > 1000 {print}'

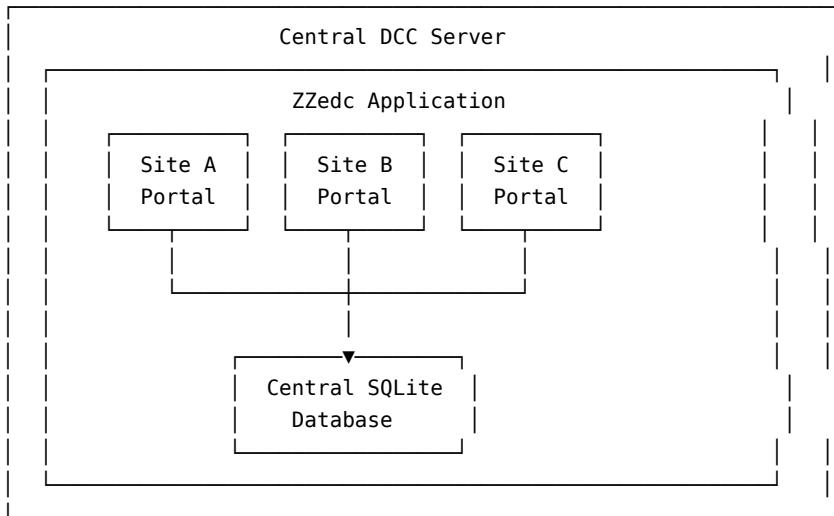
# Authentication failures
grep "Authentication failed" /var/log/zzedc/security.log |
  awk '{print $NF}' | sort | uniq -c | sort -rn

```

5. Multi-Site Coordination

5.1 Multi-Site Architecture

Centralized Database Model



5.2 Site Configuration

Site-Specific Settings

```

# /opt/zzedc/config/sites.yml

sites:
  site_001:
    name: "University Medical Center"
    code: "UMC"
    timezone: "America/New_York"
    contact: "pi@umc.edu"
    ip_whitelist:
      - "192.168.1.0/24"
      - "10.0.0.0/8"

  site_002:
    name: "Regional Hospital"
    code: "RGH"
    timezone: "America/Chicago"
    contact: "coordinator@rgh.org"
    ip_whitelist:
      - "172.16.0.0/16"

  site_003:
    name: "Research Clinic"
    code: "RCL"
    timezone: "America/Los_Angeles"
    contact: "admin@rcl.edu"
    ip_whitelist:
      - "203.0.113.0/24"

global:
  session_timeout_minutes: 30
  max_concurrent_users_per_site: 10
  data_lock_timeout_seconds: 300
  
```

5.3 Site Access Control

```
# Site-based access control
validate_site_access <- function(user_id, site_id, action, session_info) {
  user <- get_user_by_id(user_id)

  # Check site assignment
  if (!site_id %in% user$assigned_sites) {
    log_security_event("unauthorized_site_access", user_id, site_id)
    return(list(allowed = FALSE, reason = "User not assigned to site"))
  }

  # Check IP whitelist
  if (!is_ip_allowed(session_info$ip, get_site_whitelist(site_id))) {
    log_security_event("ip_not_whitelisted", user_id, site_id)
    return(list(allowed = FALSE, reason = "Access from unauthorized IP"))
  }

  # Check role permissions for action
  if (!has_permission(user$role, action)) {
    log_security_event("permission_denied", user_id, action)
    return(list(allowed = FALSE, reason = "Insufficient permissions"))
  }

  list(allowed = TRUE)
}
```

5.4 Cross-Site Data Queries

```
# Query data across sites with appropriate filtering
query_cross_site_data <- function(user_id, query_params) {
  user <- get_user_by_id(user_id)

  # DCC staff can see all sites
  if (user$role %in% c("dcc_admin", "biostatistician")) {
    site_filter <- NULL
  } else {
    # Others see only their assigned sites
    site_filter <- user$assigned_sites
  }

  # Build query with site filter
  base_query <- "SELECT * FROM form_data fd
    INNER JOIN subjects s ON fd.subject_id = s.id
    WHERE 1=1"

  if (!is.null(site_filter)) {
    base_query <- paste(base_query,
      sprintf("AND s.site_id IN (%s)",
        paste(shQuote(site_filter), collapse = ",")))
  }

  # Add user-specified filters
  if (!is.null(query_params$visit_id)) {
    base_query <- paste(base_query,
```

```

        sprintf("AND fd.visit_id = %d", query_params$visit_id))
    }

DBI::dbGetQuery(db_con, base_query)
}

5.5 Site Synchronization Monitoring

# Monitor site activity and data submission
get_site_status_summary <- function() {
  query <- "
    SELECT
      s.site_id,
      sites.name as site_name,
      COUNT(DISTINCT s.id) as enrolled_subjects,
      COUNT(DISTINCT CASE WHEN s.status = 'active' THEN s.id END) as active_subjects,
      MAX(fd.entry_timestamp) as last_data_entry,
      COUNT(CASE WHEN q.status = 'open' THEN 1 END) as open_queries
    FROM subjects s
    LEFT JOIN form_data fd ON s.id = fd.subject_id
    LEFT JOIN data_queries q ON s.id = q.subject_id
    LEFT JOIN site_config sites ON s.site_id = sites.site_id
    GROUP BY s.site_id, sites.name
    ORDER BY sites.name
  "

  status <- DBI::dbGetQuery(db_con, query)

  # Flag sites with no recent activity
  status$days_since_entry <- as.numeric(
    difftime(Sys.time(), status$last_data_entry, units = "days")
  )
  status$activity_flag <- ifelse(
    status$days_since_entry > 7, "warning",
    ifelse(status$days_since_entry > 14, "critical", "ok")
  )

  status
}

```

6. Security Operations

6.1 Security Monitoring

Failed Login Monitoring

```

# Monitor and respond to failed login attempts
monitor_failed_logins <- function(window_minutes = 15, threshold = 5) {
  cutoff <- Sys.time() - (window_minutes * 60)

  failed_attempts <- DBI::dbGetQuery(db_con, sprintf("
    SELECT
      username,
      ip_address,

```

```

        COUNT(*) as attempt_count,
        MAX(timestamp) as last_attempt
    FROM security_events
    WHERE event_type = 'failed_login'
        AND timestamp > '%s'
    GROUP BY username, ip_address
    HAVING COUNT(*) >= %d
", format(cutoff, "%Y-%m-%d %H:%M:%S"), threshold))

if (nrow(failed_attempts) > 0) {
  for (i in seq_len(nrow(failed_attempts))) {
    row <- failed_attempts[i, ]

    # Lock account if threshold exceeded
    if (row$attempt_count >= threshold * 2) {
      lock_user_account(row$username,
        reason = sprintf("Excessive failed logins from %s", row$ip_address),
        duration_minutes = 30)
    }

    # Log security alert
    log_security_event("brute_force_detected", list(
      username = row$username,
      ip_address = row$ip_address,
      attempts = row$attempt_count
    ))
  }
}

failed_attempts
}

```

6.2 Session Management

```

# Session security configuration
session_security_config <- list(
  # Session timeout (30 minutes of inactivity)
  timeout_minutes = 30,

  # Absolute session lifetime (8 hours)
  max_lifetime_hours = 8,

  # Concurrent session limit per user
  max_concurrent_sessions = 2,

  # Session regeneration interval
  regenerate_id_minutes = 15,

  # Require re-authentication for sensitive operations
  reauth_timeout_minutes = 5
)

# Enforce session limits
enforce_session_limits <- function(user_id) {

```

```

active_sessions <- get_user_active_sessions(user_id)

if (length(active_sessions) >= session_security_config$max_concurrent_sessions) {
  # Terminate oldest session
  oldest <- active_sessions[order(active_sessions$created_at)][1]
  terminate_session(oldest$session_id,
    reason = "Concurrent session limit exceeded")
}
}

```

6.3 Audit Log Protection

```

# Ensure audit log integrity
verify_audit_integrity <- function(start_date = NULL, end_date = NULL) {
  where_clause <- "WHERE 1=1"

  if (!is.null(start_date)) {
    where_clause <- paste(where_clause,
      sprintf("AND timestamp >= '%s'", start_date))
  }
  if (!is.null(end_date)) {
    where_clause <- paste(where_clause,
      sprintf("AND timestamp <= '%s'", end_date))
  }

  # Verify hash chain integrity
  audit_records <- DBI::dbGetQuery(db_con, sprintf(
    "SELECT id, timestamp, event_type, details, previous_hash, record_hash
     FROM audit_trail
     %s
     ORDER BY id
   ", where_clause))

  if (nrow(audit_records) == 0) {
    return(list(valid = TRUE, message = "No records in range"))
  }

  # Verify each record's hash
  broken_chain <- NULL
  for (i in seq_len(nrow(audit_records))) {
    record <- audit_records[i, ]

    expected_hash <- compute_audit_hash(
      record$timestamp,
      record$event_type,
      record$details,
      record$previous_hash
    )

    if (expected_hash != record$record_hash) {
      broken_chain <- c(broken_chain, record$id)
    }

    # Verify chain link
  }
}

```

```

    if (i > 1 && record$previous_hash != audit_records$record_hash[i - 1]) {
      broken_chain <- c(broken_chain, record$id)
    }
  }

  if (length(broken_chain) > 0) {
    log_security_event("audit_integrityViolation", list(
      affected_records = broken_chain
    ))
  }

  return(list(
    valid = FALSE,
    message = "Audit chain integrity violation detected",
    affected_records = broken_chain
  ))
}

list(valid = TRUE, message = "Audit integrity verified", records_checked = nrow(audit_records))
}

```

7. Upgrade and Patch Management

7.1 Version Upgrade Procedure

Pre-Upgrade Checklist

```
## Pre-Upgrade Checklist
```

- [] Review release notes for breaking changes
- [] Verify backup completed within last 24 hours
- [] Test upgrade in staging environment
- [] Notify users of maintenance window
- [] Confirm rollback procedure is ready
- [] Document current version and configuration
- [] Verify disk space for upgrade (>= 2GB free)
- [] Check R package compatibility matrix

Upgrade Script

```

#!/bin/bash
# /opt/zzedc/scripts/upgrade.sh

set -e

NEW_VERSION=$1
BACKUP_DIR="/var/backups/zzedc/pre-upgrade"
INSTALL_DIR="/opt/zzedc"

if [ -z "$NEW_VERSION" ]; then
  echo "Usage: $0 <version>"
  exit 1
fi

echo "Upgrading ZZedc to version $NEW_VERSION"

```

```

# 1. Create pre-upgrade backup
echo "Creating pre-upgrade backup..."
mkdir -p "$BACKUP_DIR"
/opt/zzedc/scripts/daily_backup.sh
cp -r "$INSTALL_DIR" "$BACKUP_DIR/zzedc_$(date +%Y%m%d)"

# 2. Stop service
echo "Stopping ZZedc service..."
sudo systemctl stop zzedc

# 3. Download new version
echo "Downloading version $NEW_VERSION..."
cd /tmp
wget "https://github.com/org/zzedc/archive/v${NEW_VERSION}.tar.gz"
tar -xzf "v${NEW_VERSION}.tar.gz"

# 4. Install new version
echo "Installing new version..."
cp -r "zzedc-${NEW_VERSION}/*" "$INSTALL_DIR/"

# 5. Update R packages
echo "Updating R dependencies..."
cd "$INSTALL_DIR"
Rscript -e "renv::restore()"

# 6. Run database migrations
echo "Running database migrations..."
Rscript -e "zzedc::run_migrations()"

# 7. Verify installation
echo "Verifying installation..."
Rscript -e "zzedc::verify_installation()"

# 8. Start service
echo "Starting ZZedc service..."
sudo systemctl start zzedc

# 9. Health check
sleep 10
if curl -s http://localhost:3838/health | grep -q "ok"; then
    echo "Upgrade to version $NEW_VERSION completed successfully"
else
    echo "Health check failed. Initiating rollback..."
    /opt/zzedc/scripts/rollback.sh
    exit 1
fi

```

7.2 Rollback Procedure

```

#!/bin/bash
# /opt/zzedc/scripts/rollback.sh

BACKUP_DIR="/var/backups/zzedc/pre-upgrade"

```

```

INSTALL_DIR="/opt/zzedc"

# Find most recent backup
LATEST_BACKUP=$(ls -td "$BACKUP_DIR"/zzedc_* | head -1)

if [ -z "$LATEST_BACKUP" ]; then
    echo "No backup found for rollback"
    exit 1
fi

echo "Rolling back to: $LATEST_BACKUP"

# Stop service
sudo systemctl stop zzedc

# Restore application
rm -rf "$INSTALL_DIR"
cp -r "$LATEST_BACKUP" "$INSTALL_DIR"

# Restore database
LATEST_DB_BACKUP=$(ls -t "$BACKUP_DIR"/*.db 2>/dev/null | head -1)
if [ -n "$LATEST_DB_BACKUP" ]; then
    cp "$LATEST_DB_BACKUP" /var/data/zzedc/study.db
fi

# Start service
sudo systemctl start zzedc

echo "Rollback completed"

```

7.3 Security Patch Policy

Security Patch Response Times

Severity	Description	Response Time	Examples
Critical	Active exploitation, data breach risk	24 hours	RCE, SQL injection, auth bypass
High	Significant vulnerability, no active exploit	7 days	XSS, CSRF, privilege escalation
Medium	Limited impact or difficult to exploit	30 days	Information disclosure, DoS
Low	Minimal security impact	Next release	UI issues, minor info leaks

Patch Verification

After applying security patches:

1. Run full test suite: `Rscript -e "devtools::test()"`
 2. Verify affected functionality manually
 3. Check security scanner: `Rscript -e "zzedc::security_scan()"`
 4. Review audit logs for anomalies
 5. Document patch application in change log
-

Appendix A: Command Reference

System Management

```
# Service control
sudo systemctl start zzedc
sudo systemctl stop zzedc
sudo systemctl restart zzedc
sudo systemctl status zzedc

# Log viewing
sudo journalctl -u zzedc -f
tail -f /var/log/zzedc/app.log
tail -f /var/log/zzedc/error.log

# Database management
sqlite3 /var/data/zzedc/study.db ".tables"
sqlite3 /var/data/zzedc/study.db "PRAGMA integrity_check"
sqlite3 /var/data/zzedc/study.db "VACUUM"
```

R Console Commands

```
# Database health
zzedc::check_database_health("/var/data/zzedc/study.db")

# User management
zzedc::list_users()
zzedc::reset_user_password("username")
zzedc::unlock_user_account("username")

# Audit review
zzedc::get_recent_audit_events(n = 100)
zzedc::search_audit_trail(user = "admin", action = "delete")

# System status
zzedc::get_system_status()
zzedc::get_active_sessions()
```

Appendix B: Monitoring Dashboards

Recommended Metrics for Grafana/Prometheus

```
# Prometheus scrape configuration
scrape_configs:
  - job_name: 'zzedc'
    scrape_interval: 30s
    static_configs:
      - targets: ['localhost:3838']
    metrics_path: '/metrics'
```

Key Performance Indicators

Metric	Warning	Critical	Description
response_time_p95	>2s	>5s	95th percentile response time
error_rate	>1%	>5%	Request error rate
active_sessions	>80% capacity	>95% capacity	Concurrent users
db_query_time_avg	>500ms	>2s	Average query execution
memory_usage	>75%	>90%	System memory utilization
disk_usage	>80%	>95%	Database disk utilization

Document History

Version	Date	Author	Changes
1.0.0	December 2025	ZZedc Team	Initial release